

Trinity HPC Cluster Usage Overview

<https://github.com/mhibbs-trinity/TU-HPC-Guide>

leviosa.cs.trinity.edu

Head
Node

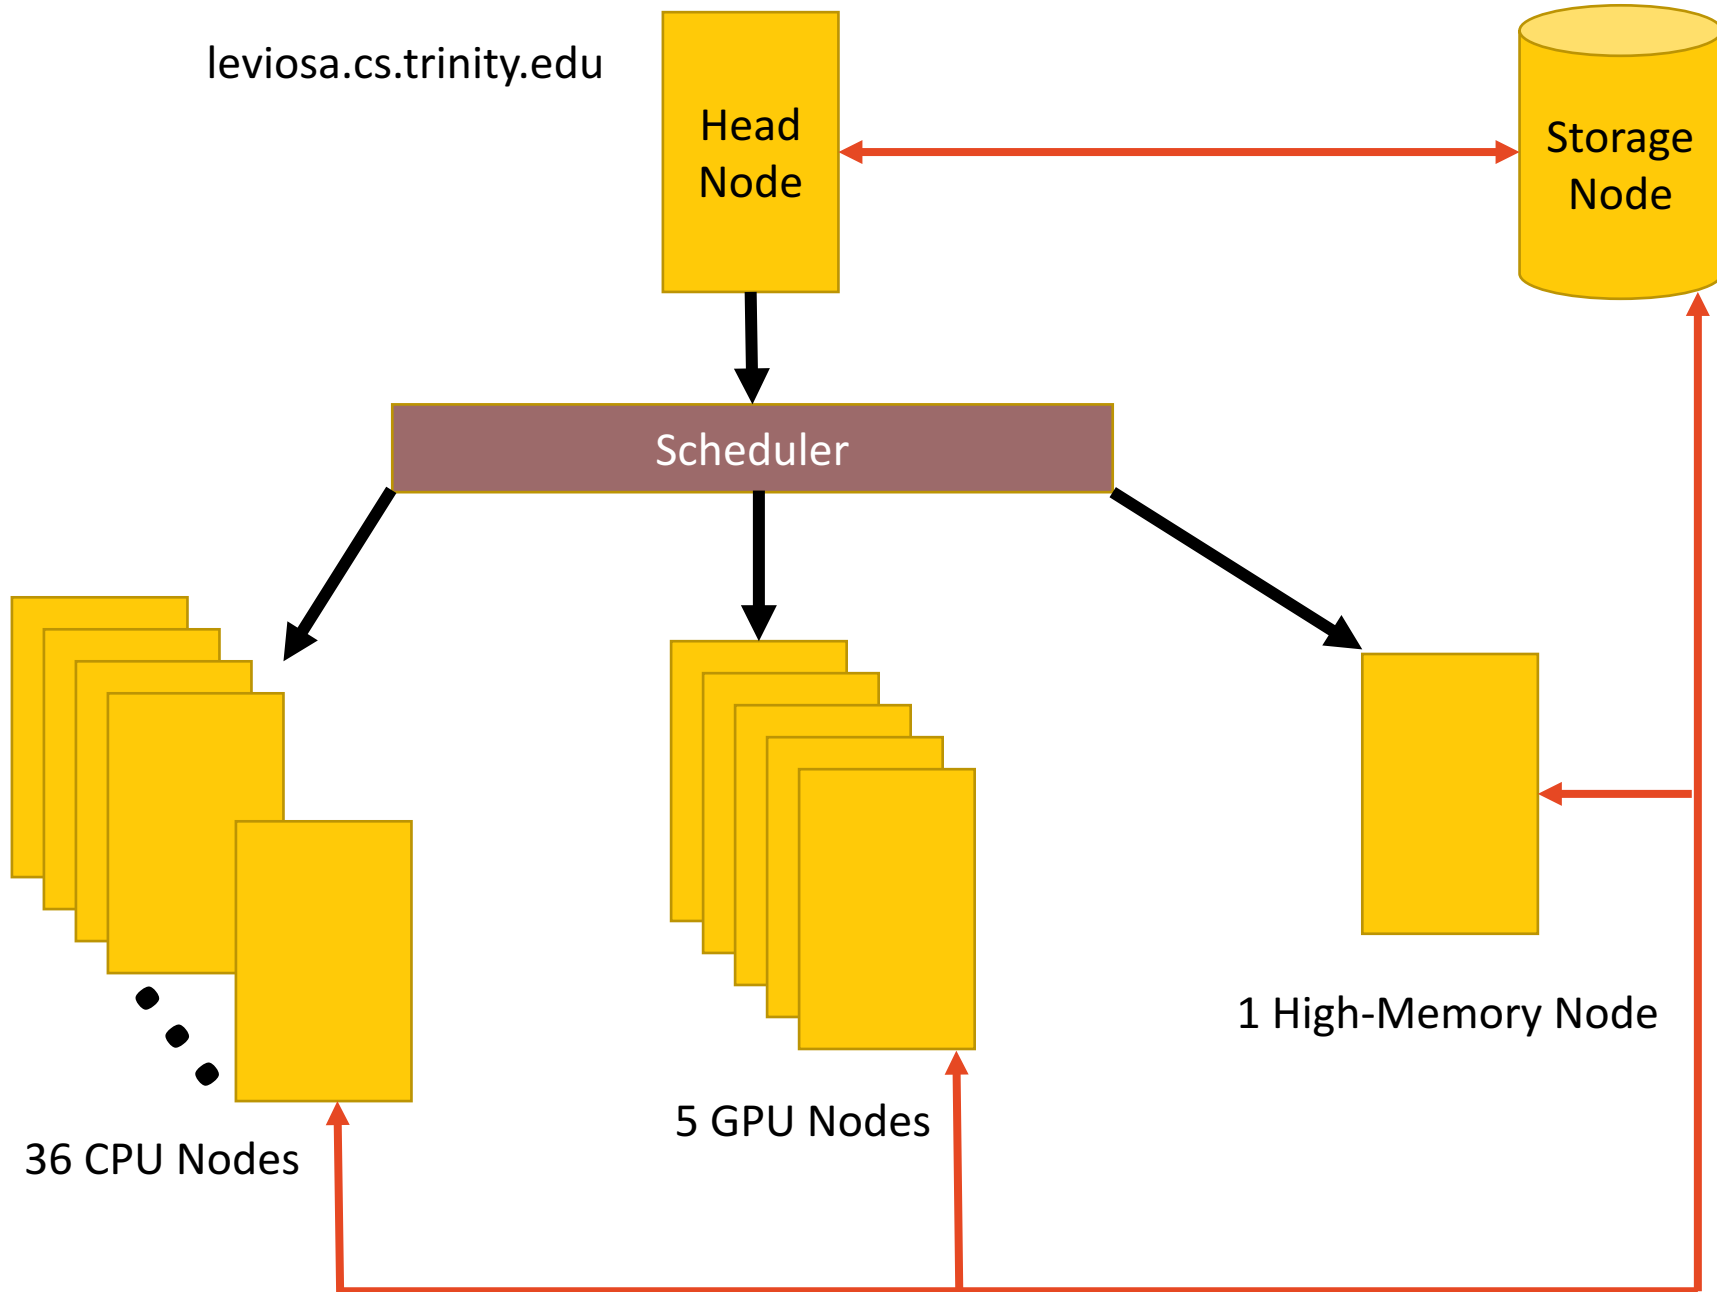
Storage
Node

Scheduler

1 High-Memory Node

5 GPU Nodes

36 CPU Nodes



leviosa.cs.trinity.edu

Dual Intel Xeon ES-2695 v4
36 cores total (18 per CPU)
2.1 GHz clock speed
64 GB of DDR4-2133 ECC RAM

Head
Node

Dual Intel Xeon ES-2620 v4
16 cores total (8 per CPU)
2.1 GHz clock speed
64 GB of DDR4-2133 ECC RAM
216 TB of raw HDD Storage
Configured as 100 TB usable RAID 6 Storage

Storage
Node

Scheduler

1 High-Memory Node

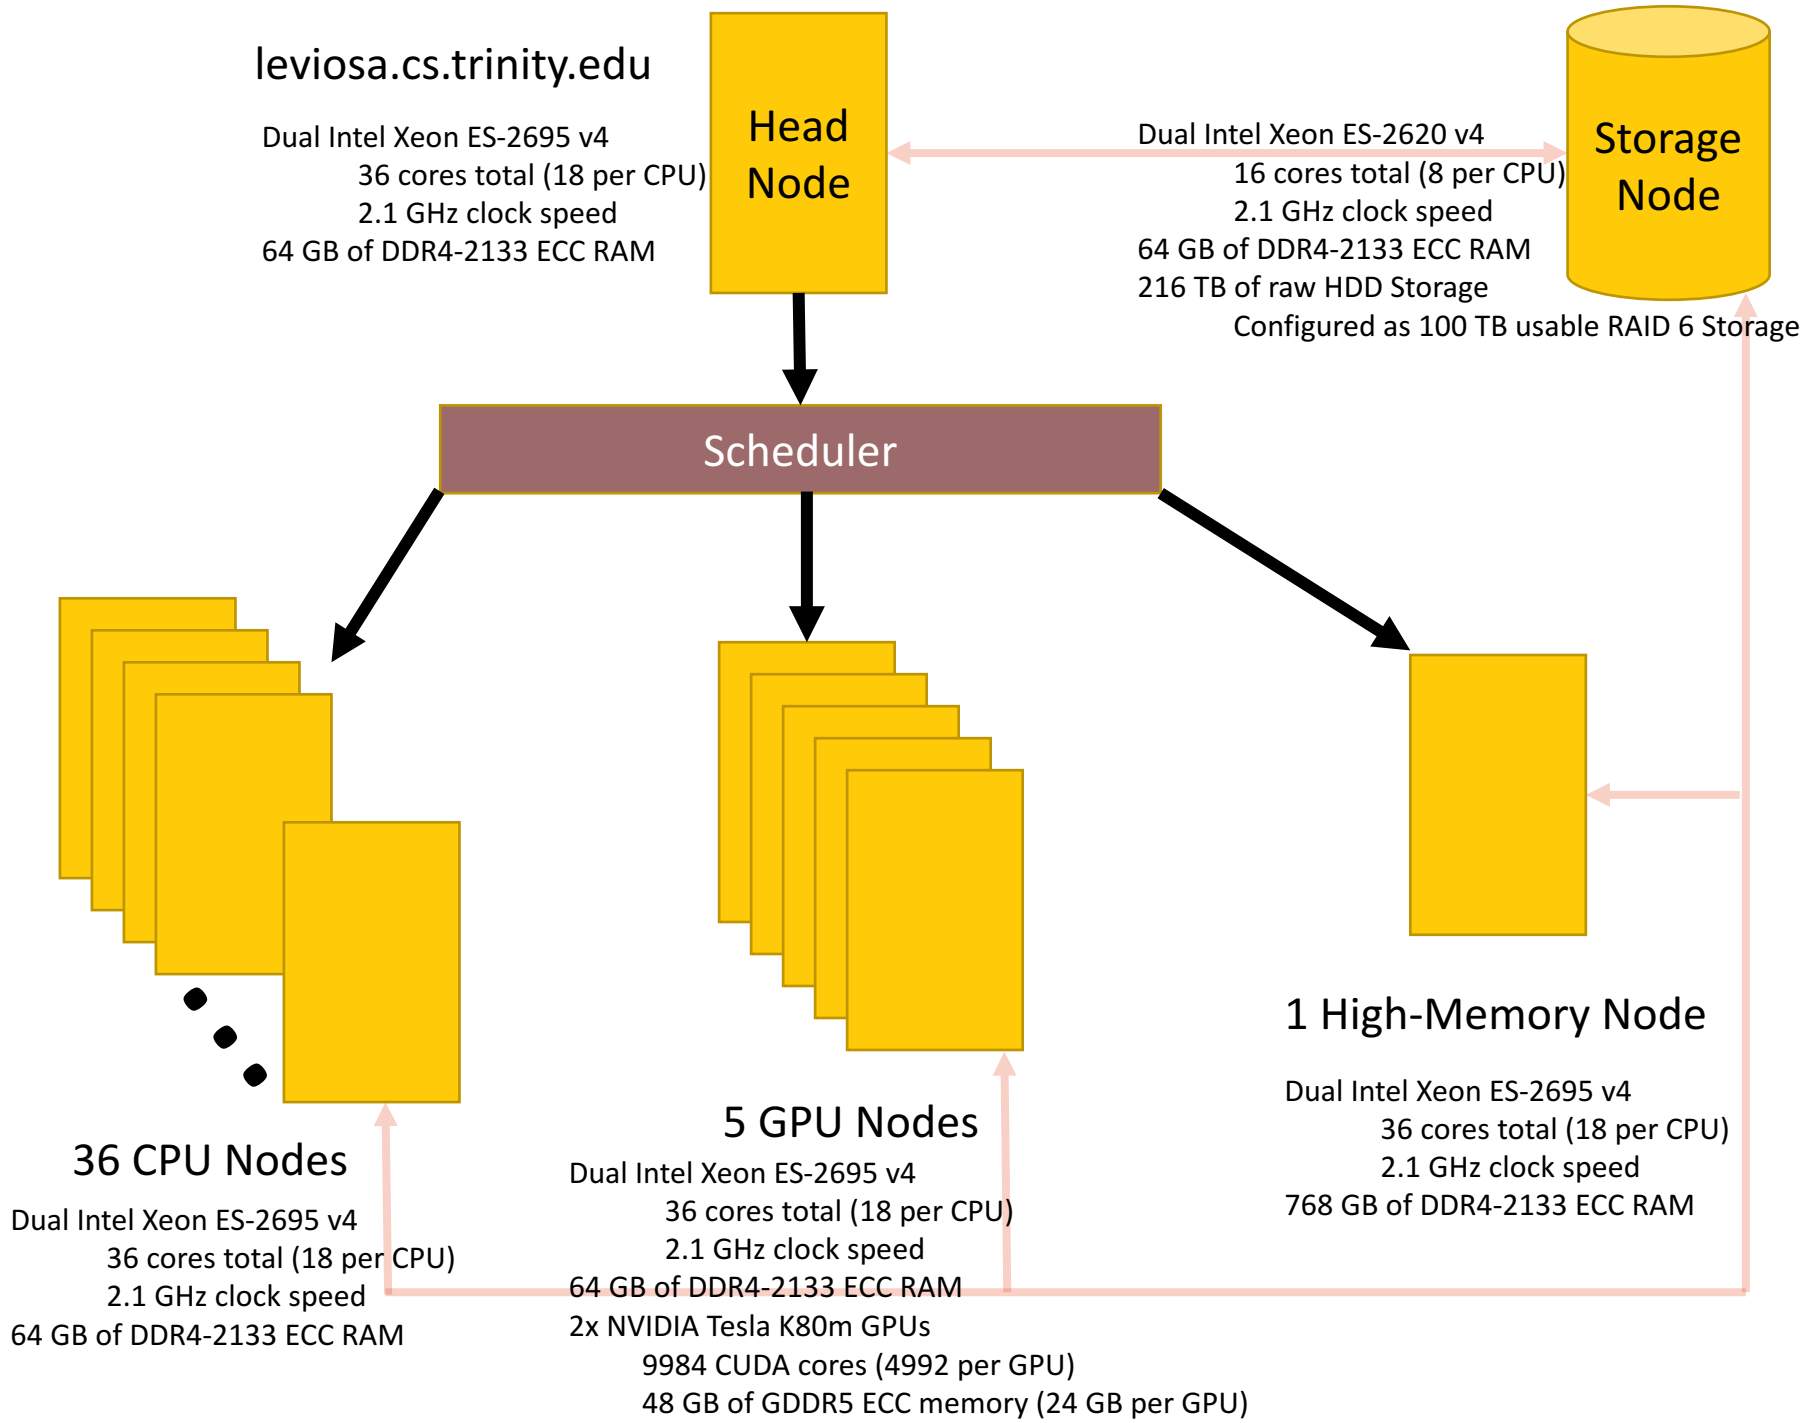
Dual Intel Xeon ES-2695 v4
36 cores total (18 per CPU)
2.1 GHz clock speed
768 GB of DDR4-2133 ECC RAM

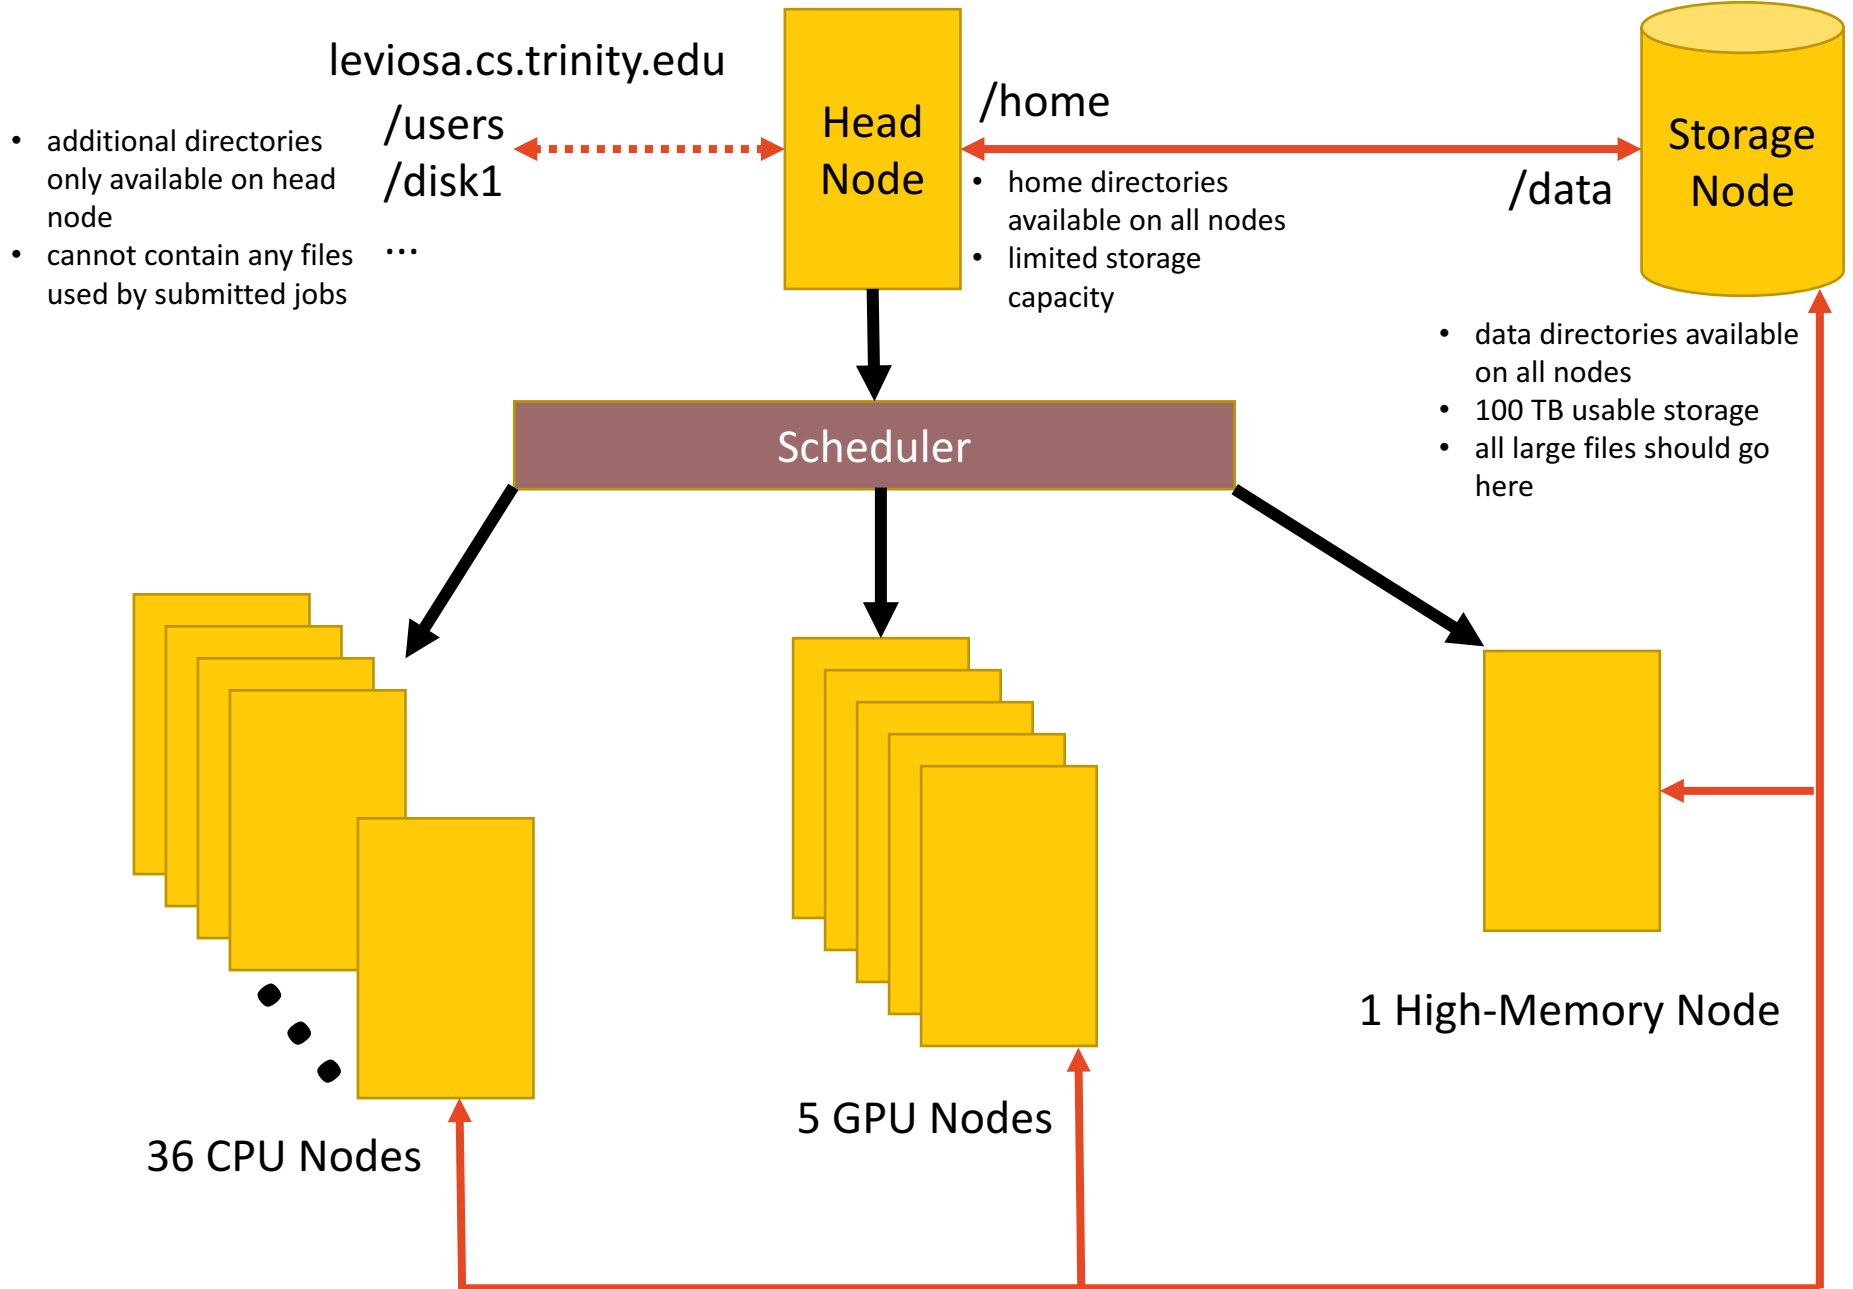
5 GPU Nodes

Dual Intel Xeon ES-2695 v4
36 cores total (18 per CPU)
2.1 GHz clock speed
64 GB of DDR4-2133 ECC RAM
2x NVIDIA Tesla K80m GPUs
9984 CUDA cores (4992 per GPU)
48 GB of GDDR5 ECC memory (24 GB per GPU)

36 CPU Nodes

Dual Intel Xeon ES-2695 v4
36 cores total (18 per CPU)
2.1 GHz clock speed
64 GB of DDR4-2133 ECC RAM





OS & Software Information

- CentOS 7.3
 - derivative of RedHat Enterprise Linux
- TORQUE scheduling manager
- Core Linux software (e.g. cpp, vim, emacs, etc)
- Contact an administrator for help installing software
 - Neal Pape (npape@trinity.edu)
 - Matt Hibbs (mhibbs@trinity.edu)

Logging On

- Accounts are separate from other Linux systems
 - Send account requests to Neal Pape (npape@trinity.edu)

```
varys:~ mhibbs$ ssh mhibbs@leviosa.cs.trinity.edu
mhibbs@leviosa.cs.trinity.edu's password:
Last login: Wed May 24 21:58:34 2017 from lumos.cs.trinity.edu
[mhibbs@leviosa ~]$ █
```

Preparing to Submit Jobs

- The head node should only be used for small tasks
 - Compiling software
 - Viewing results
 - Copying files
 - etc
- All large jobs/tasks must be submitted to the scheduler in order for them to run on remote nodes
 - Prepare a shell script file
 - Include appropriate PBS directives
 - Submit job with qsub

Preparing Script File

- Bash shell script files
 - Start with `#!/bin/sh`
- PBS directives at the beginning of file
 - Comments that tell the scheduler how many resources to request for a job
- Linux commands to execute your program of interest

Example Script File

```
#!/bin/sh
```

```
#PBS -N keras_cifar10_example
```

```
#PBS -l nodes=1:ppn=36
```

```
#PBS -l walltime=72:00:00
```

```
#PBS -M mhibbs@trinity.edu
```

```
#PBS -m abe
```

```
python /data/hibbslab/projects/keras/keras/examples/cifar10_cnn.py
```



Submitting Jobs & Job Status

```
[mhibbs@leviosa ~]$ qsub run_keras_cifar10.sh  
66.leviosa.cs.trinity.edu
```

```
[mhibbs@leviosa ~]$ qstat
```

Job ID	Name	User	Time Use	S	Queue
66.leviosa	...far10_example	mhibbs	0	R	batch

Removing/Stopping Jobs

```
[mhibbs@leviosa ~]$ qstat
```

Job ID	Name	User	Time Use	S	Queue
66.leviosa	...far10_example	mhibbs	00:06:51	R	batch

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$ qdel 66.leviosa
```

```
[mhibbs@leviosa ~]$ qstat
```

Job ID	Name	User	Time Use	S	Queue
66.leviosa	...far10_example	mhibbs	00:24:00	C	batch

```
[mhibbs@leviosa ~]$
```

```
[mhibbs@leviosa ~]$
```